



Do not Forget Hardware Prefetching While Designing Secure Cache System

MS Bi-annual Progress Seminar

Sumon Nath | sumon@cse.iitb.ac.in

Advised by **Prof. Biswabandan Panda**


Recap: Speculative execution attacks

Apple M1

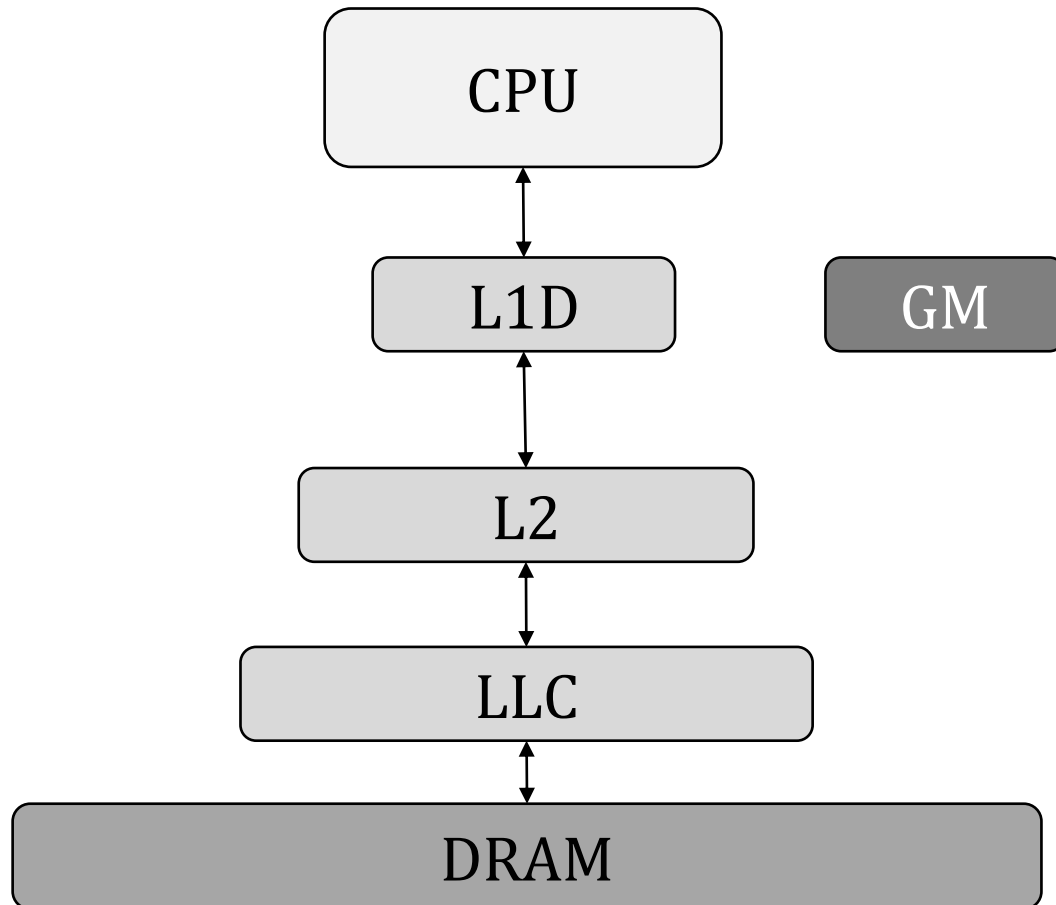
- ▶ Out-of-order execution
- ▶ Branch prediction

- ▶ Wrong path execution
- ▶ Secret data into cache

- ▶ Attacker access secret data across context switches



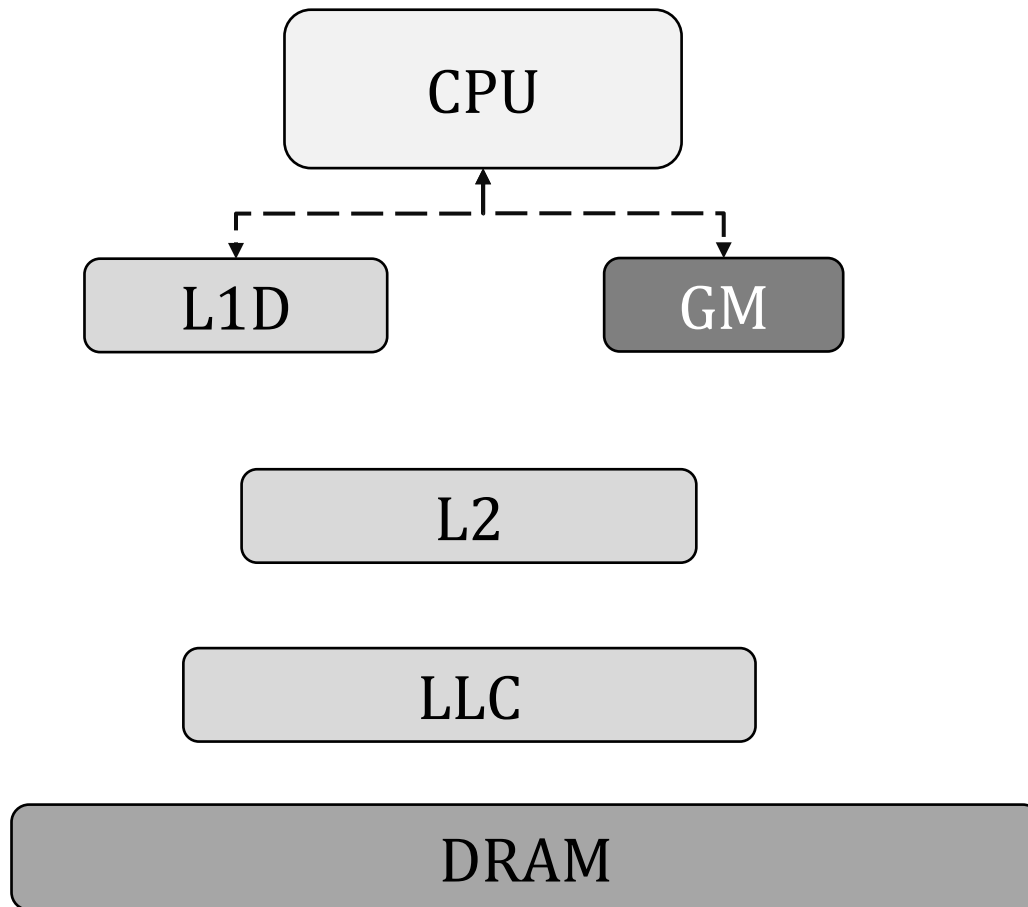
Secure Cache system: GhostMinion^[MICRO '21]




Filter cache

- Stores all speculative data
- Wiped on context switch

Secure Cache system: GhostMinion^[MICRO '21]



Secure Cache system: GhostMinion^[MICRO '21]



----- Speculative path
——— Commit Write – hit in GM
——— Commit Load – miss in GM

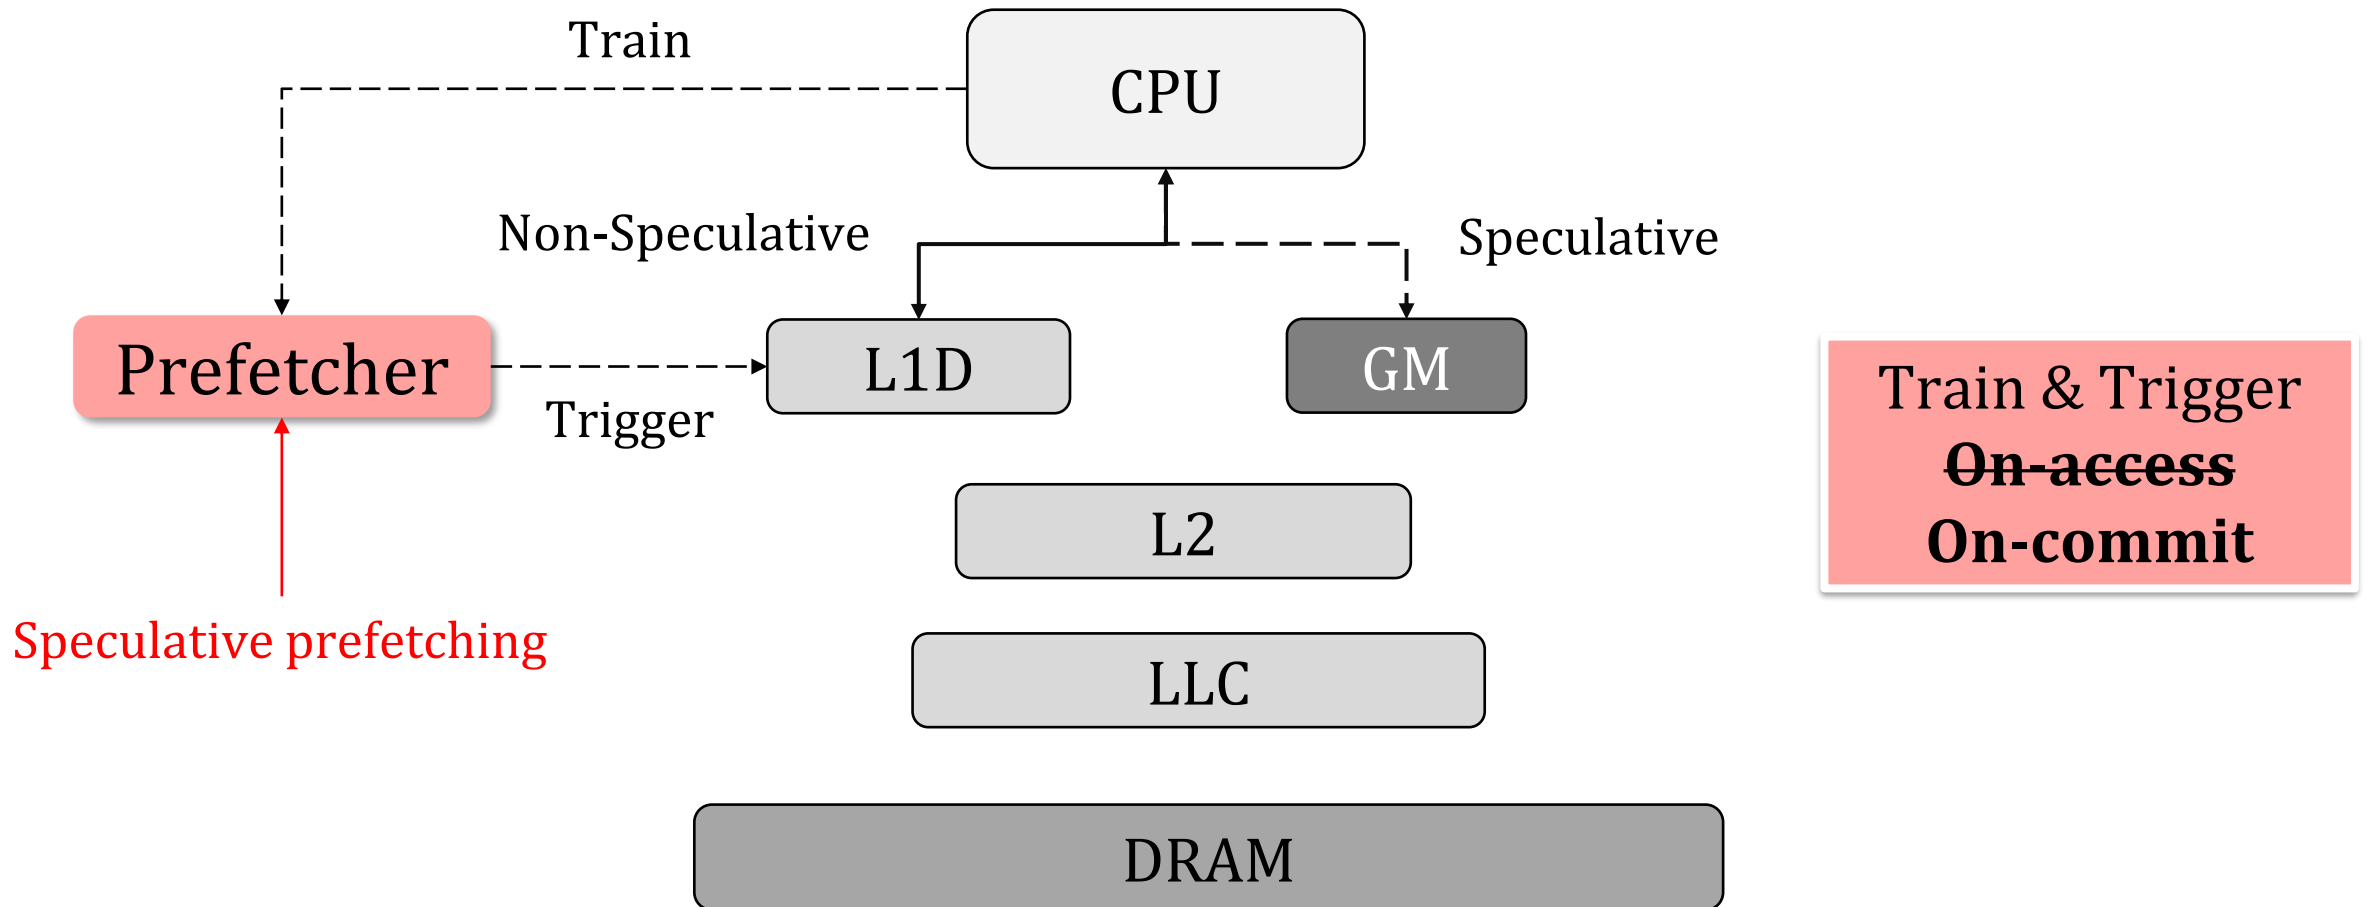
Cache state untouched

- LRU bits **not updated** on hit
- **Bypass fills** to L1/L2/LLC

Caches updated On-commit

- **Writeback** to L1-> L2 -> LLC
- **Re-fetch** to L1


What about prefetching?



Prefetcher of interest: Berti^[MICRO '22]

Accurate and **timely local delta** L1D prefetcher

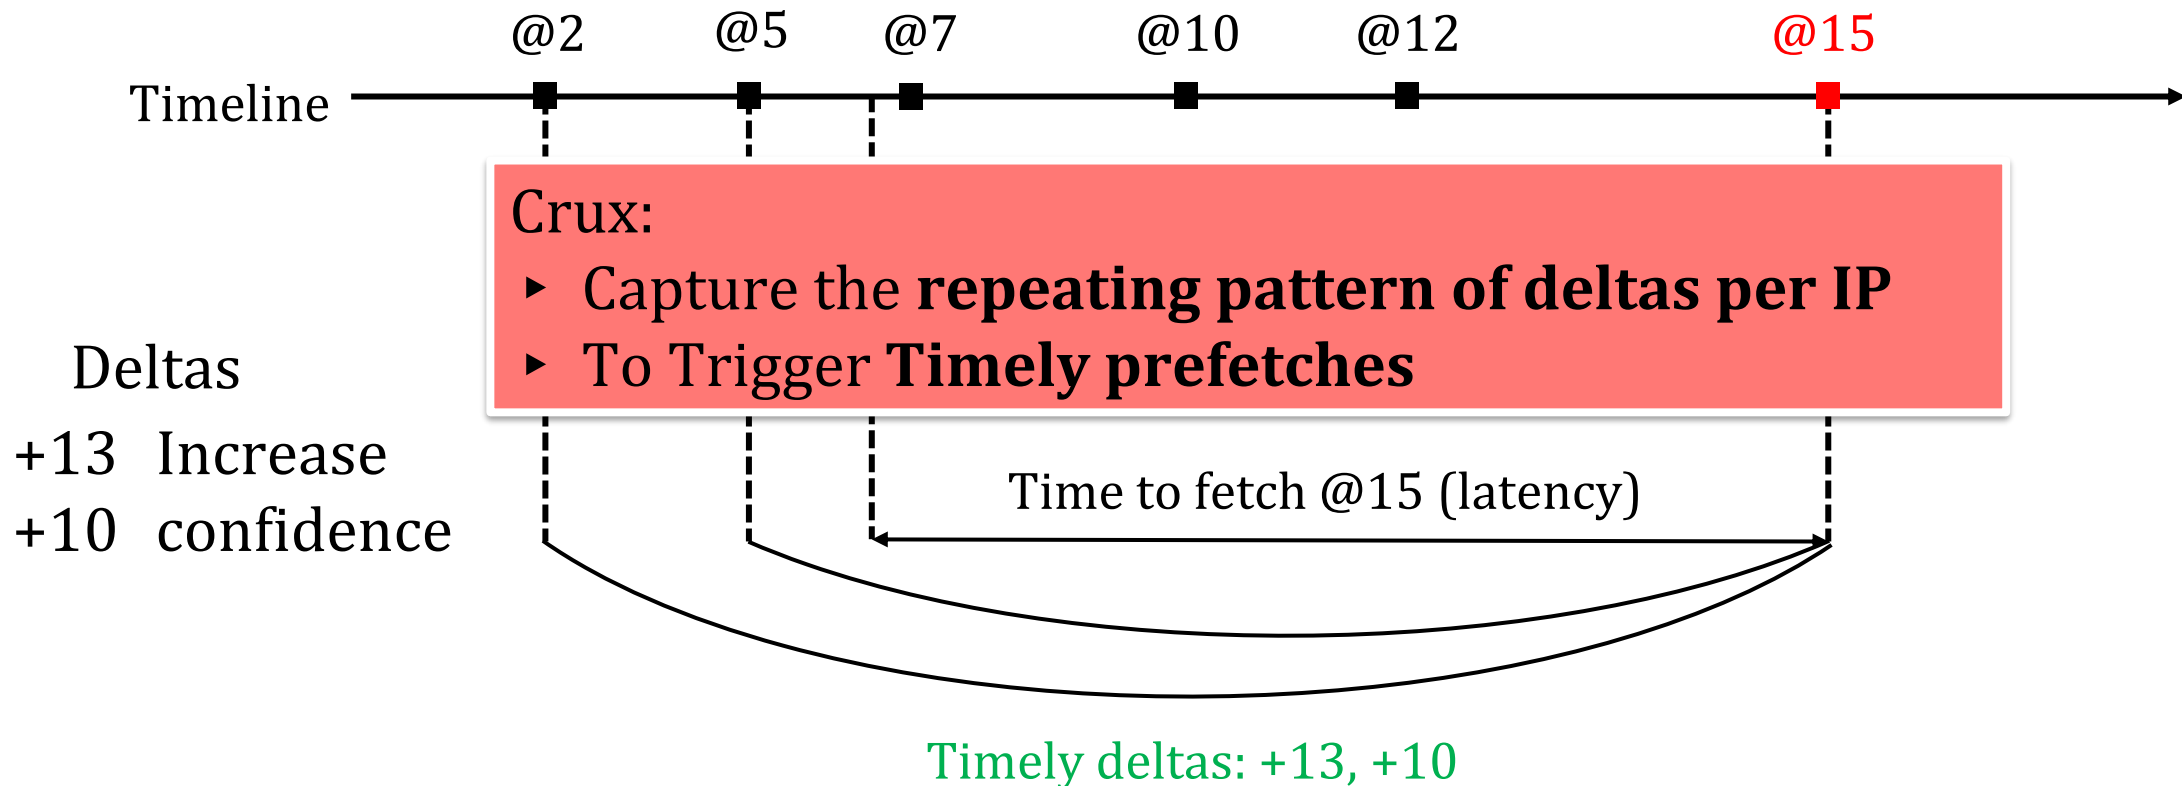
- ▶ **Delta**: Diff. b/w two cache lines
- ▶ **Local**: Per Instruction Pointer
- ▶ **Timely**: Prefetched before demand access



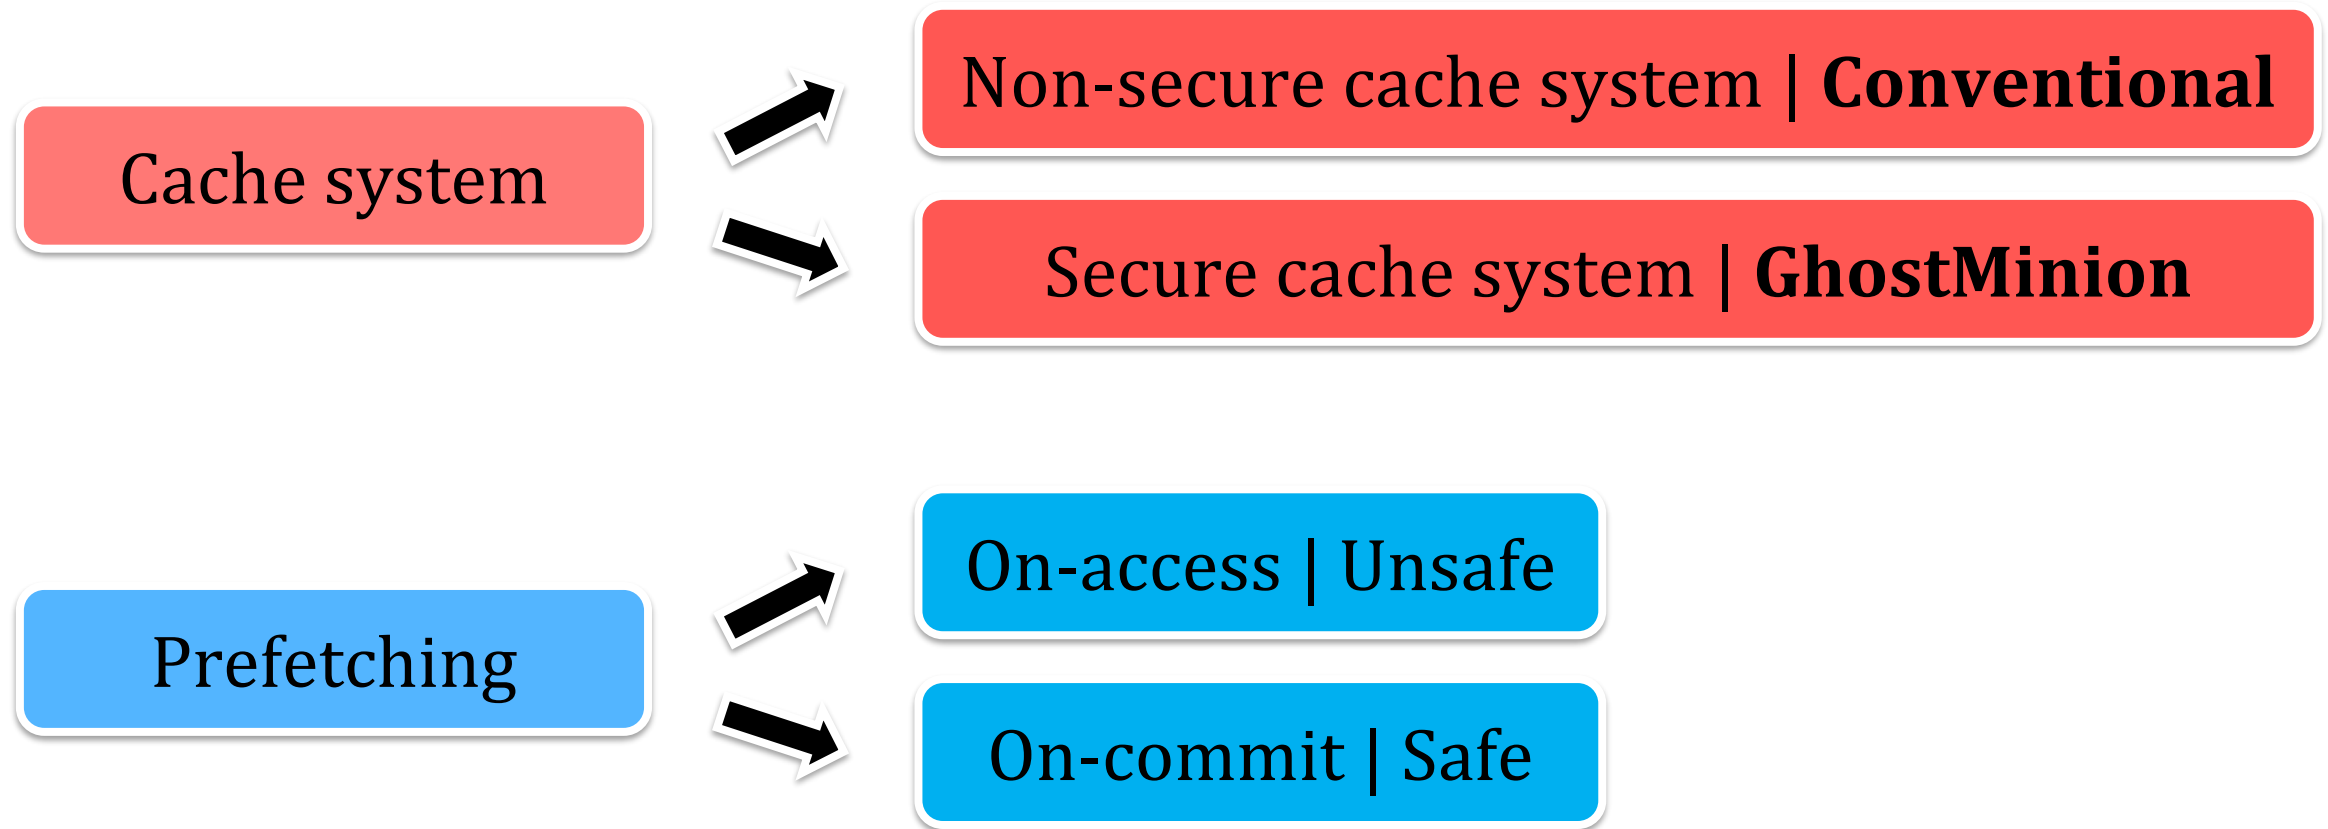
Berti^[MICRO '22] prefetcher

Accurate and **timely local delta** L1D prefetcher


- ▶ **Delta**: Diff. b/w two cache lines
- ▶ **Local**: Per Instruction Pointer
- ▶ **Timely**: Prefetched before demand access



Terminology



Impact of secure cache system & on-commit prefetching



Overall ~11 % degradation from on-access(non-secure) to on-commit (secure)

Problems & contributions

Negative impact of secure
cache system on prefetching




Secure update filter

Timeliness issues with
secure prefetching




Timely secure prefetcher

Problem 1: Impact of secure cache system




Average **Traffic** increases from **199** to **375** accesses per kilo instruction (APKI)

Increase in miss latency

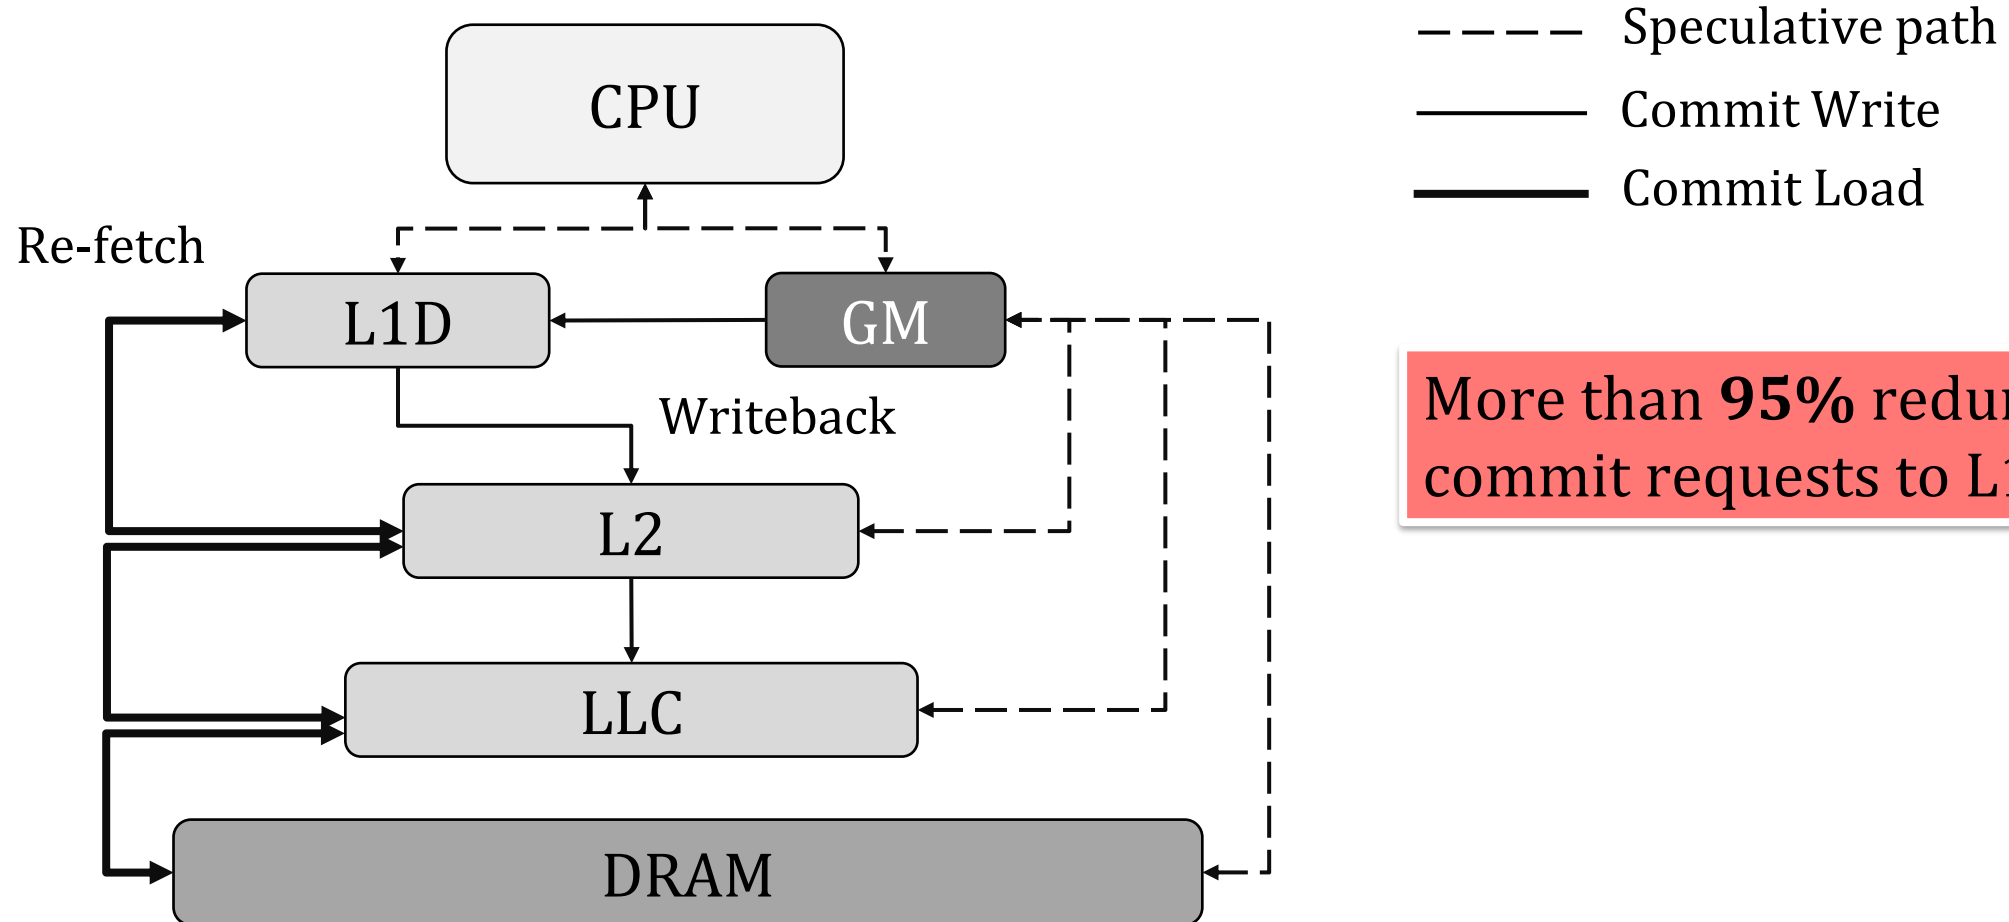


L1D load miss latency increases from 83 to 117 cycles

Increased Pressure on L1D

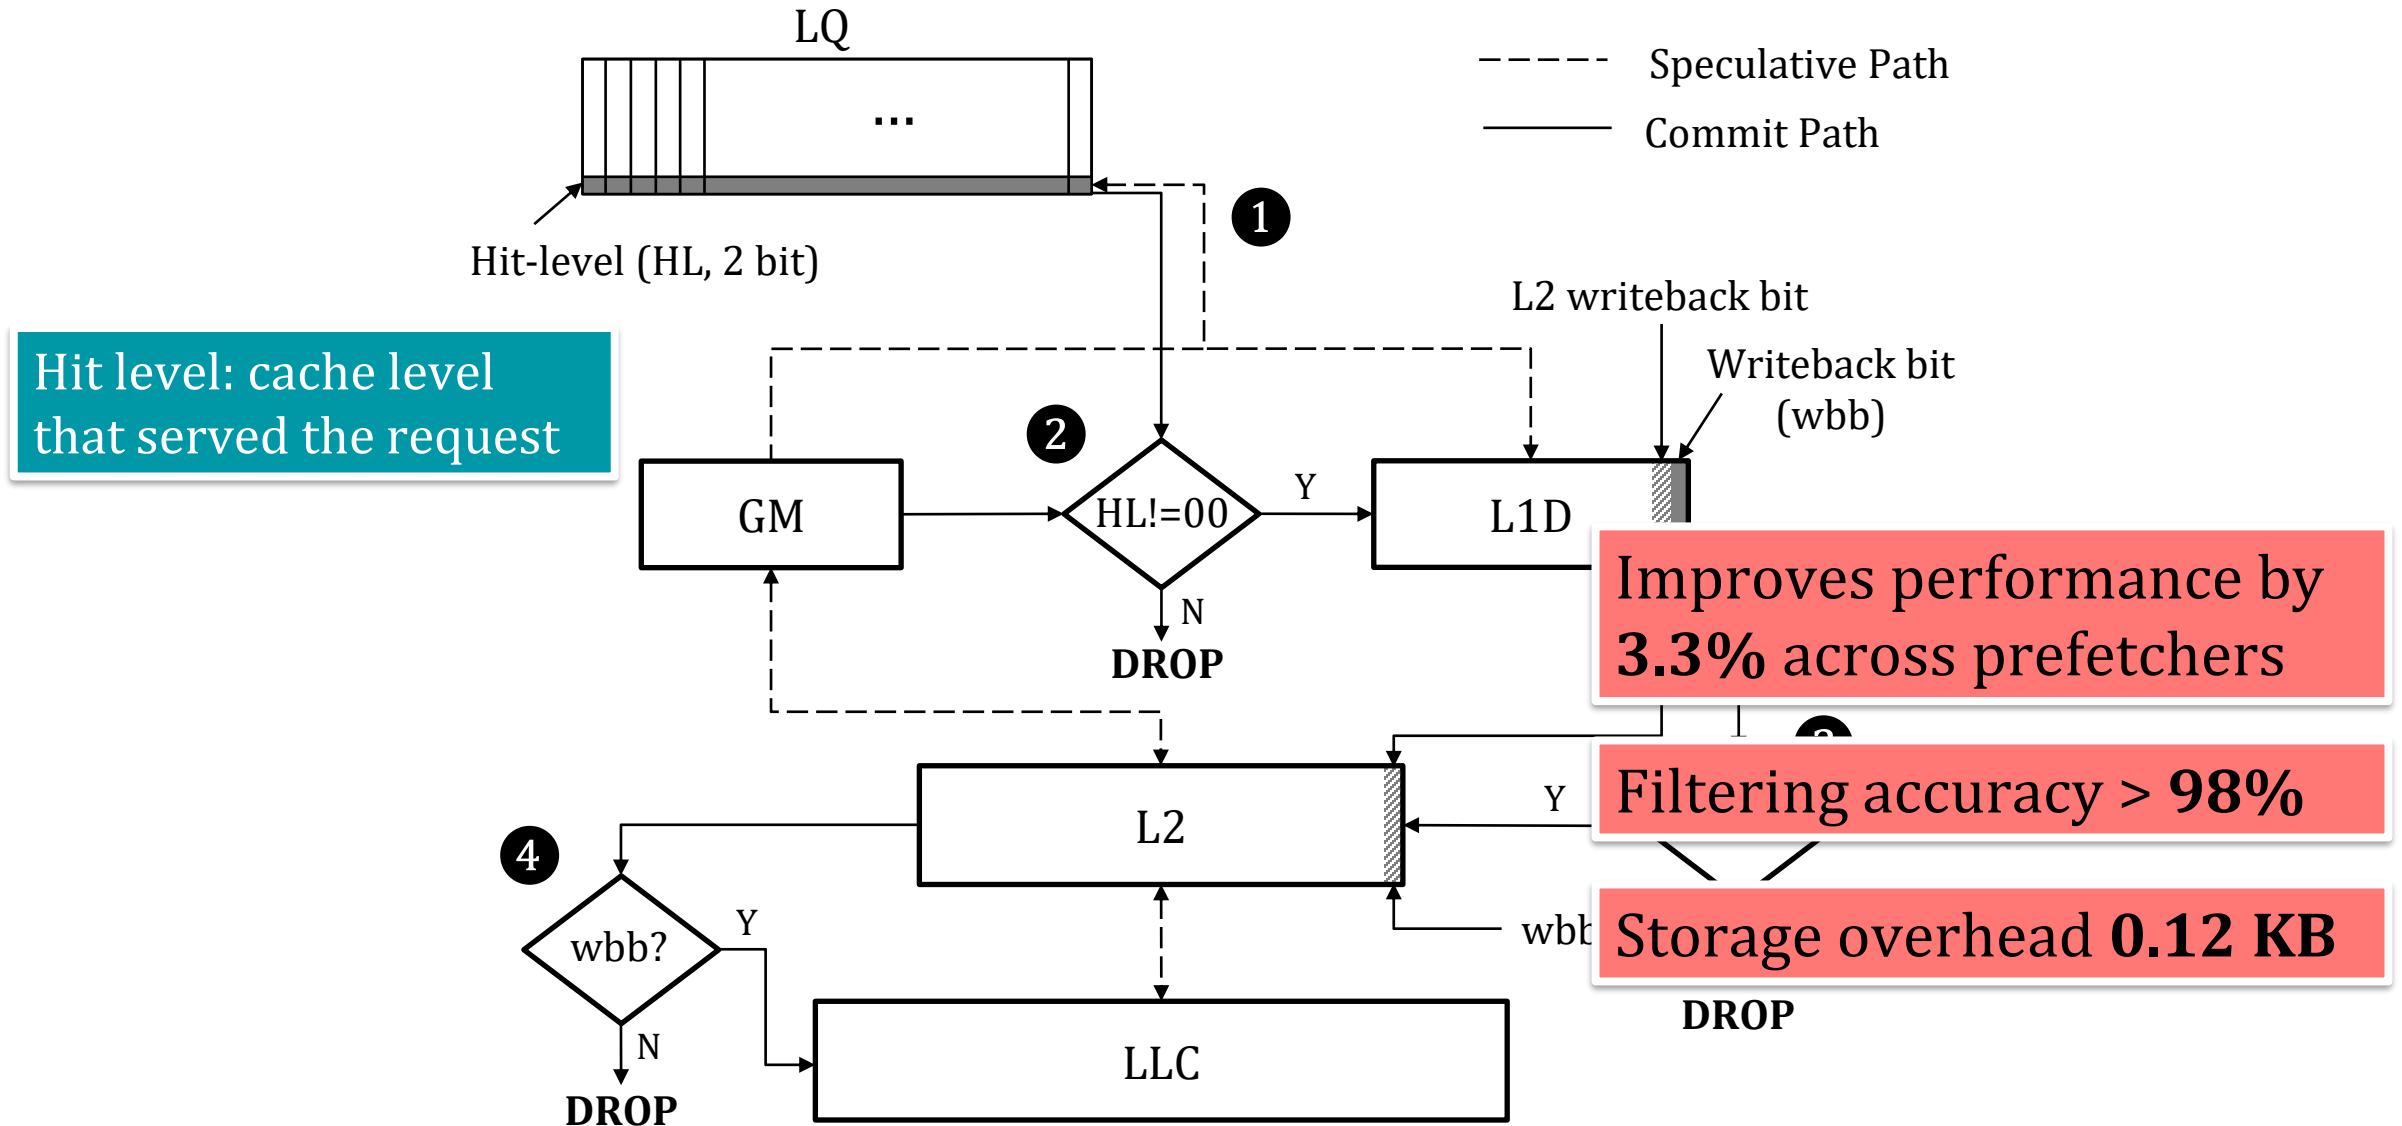


Redundant requests to cache system



More than **95%** redundant commit requests to L1D

Enhancement 1: Secure update filter(SUF)




Problem 2: Issues with Secure Prefetcher

- ▶ Berti relies on fetch latency
- ▶ Latency seen by prefetcher is mis-leading (Timeliness)
- ▶ Trigger events should be relative to on-commit events

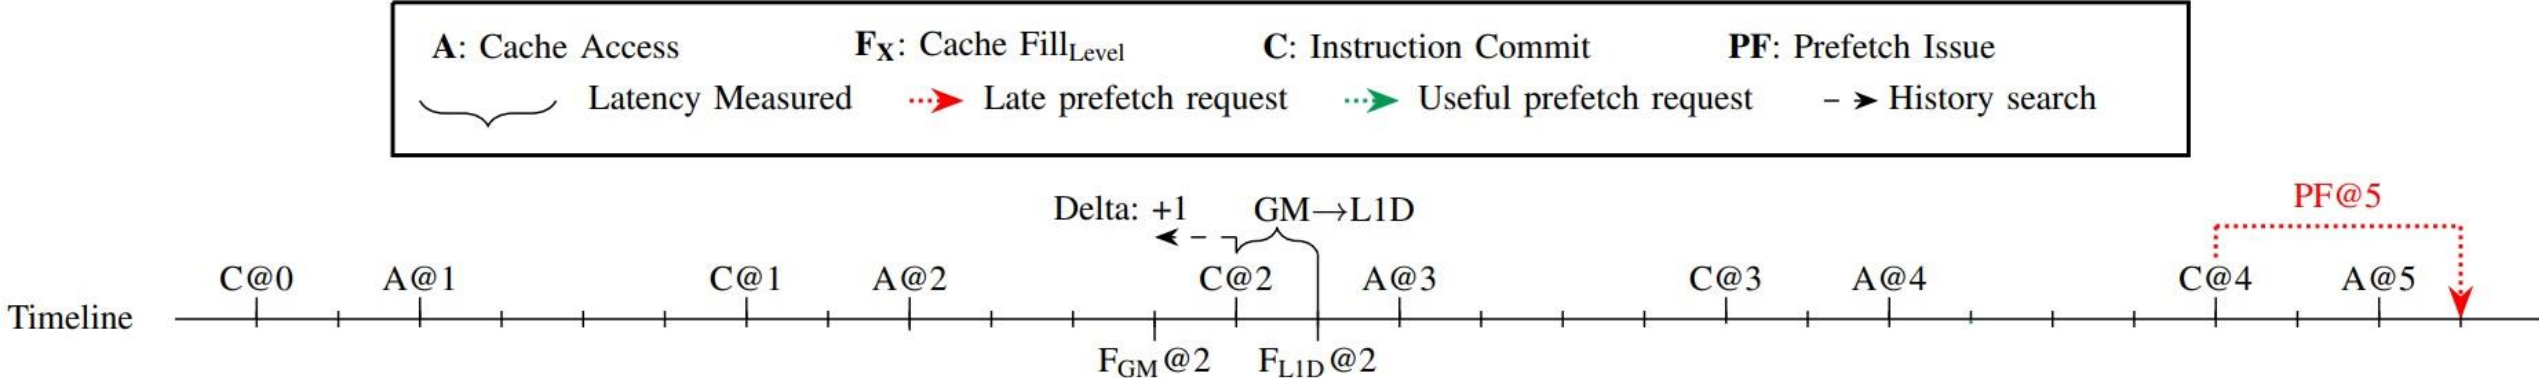
On-commit prefetching Timeliness

Timely prefetch if prefetcher trained on-access

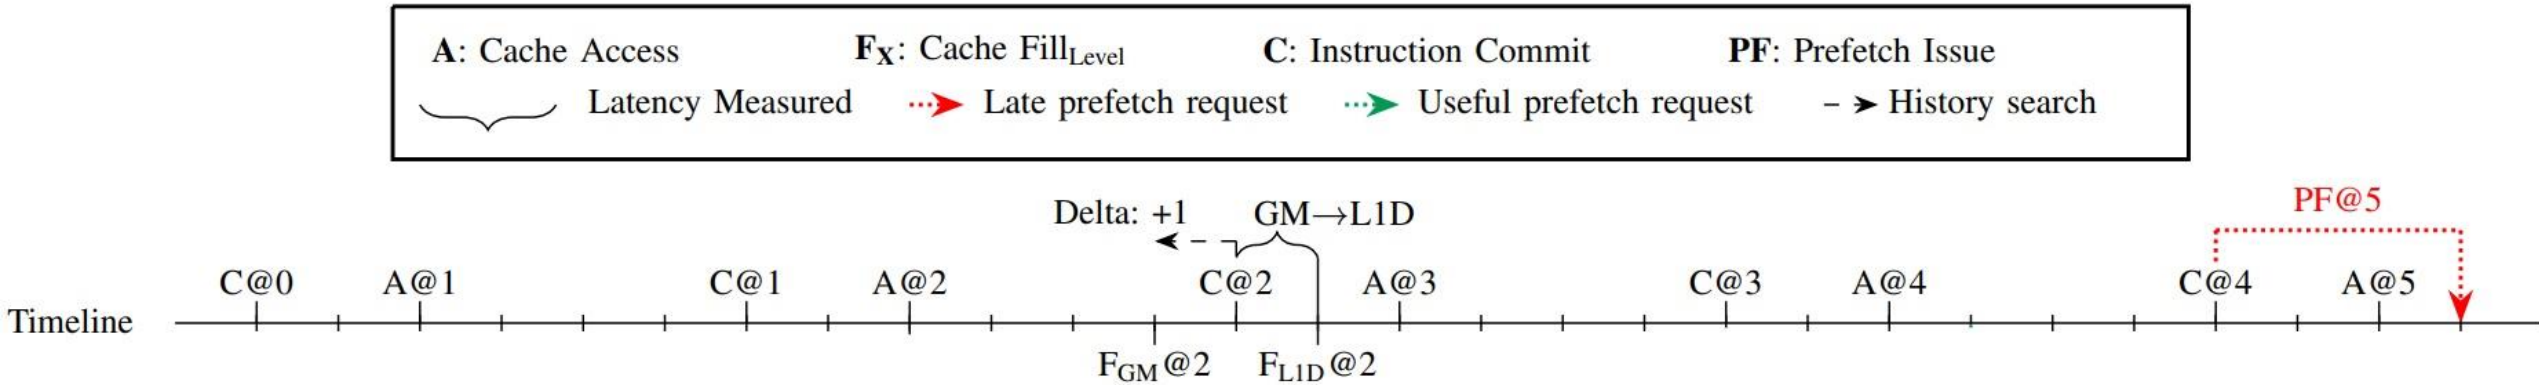


Commit Late prefetches increases overall MPKI


On-commit Berti: Measured latency problem



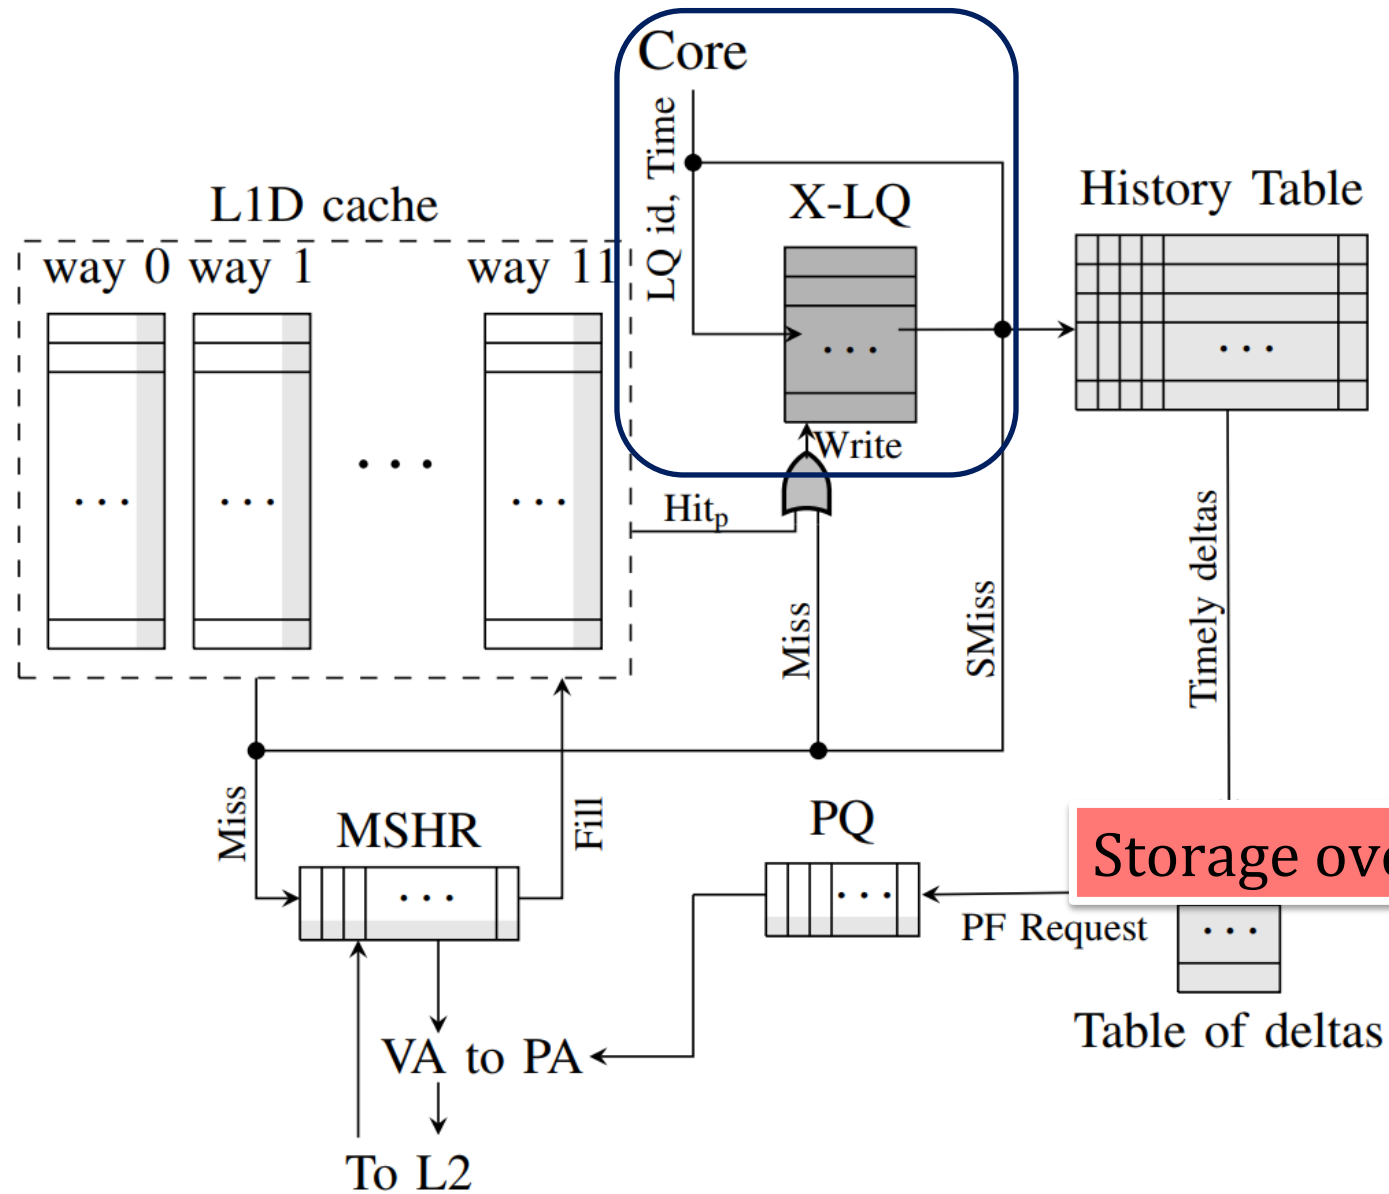
On-commit Berti: Measured latency problem



Enhancement 2: Timely Secure Berti (TSB)



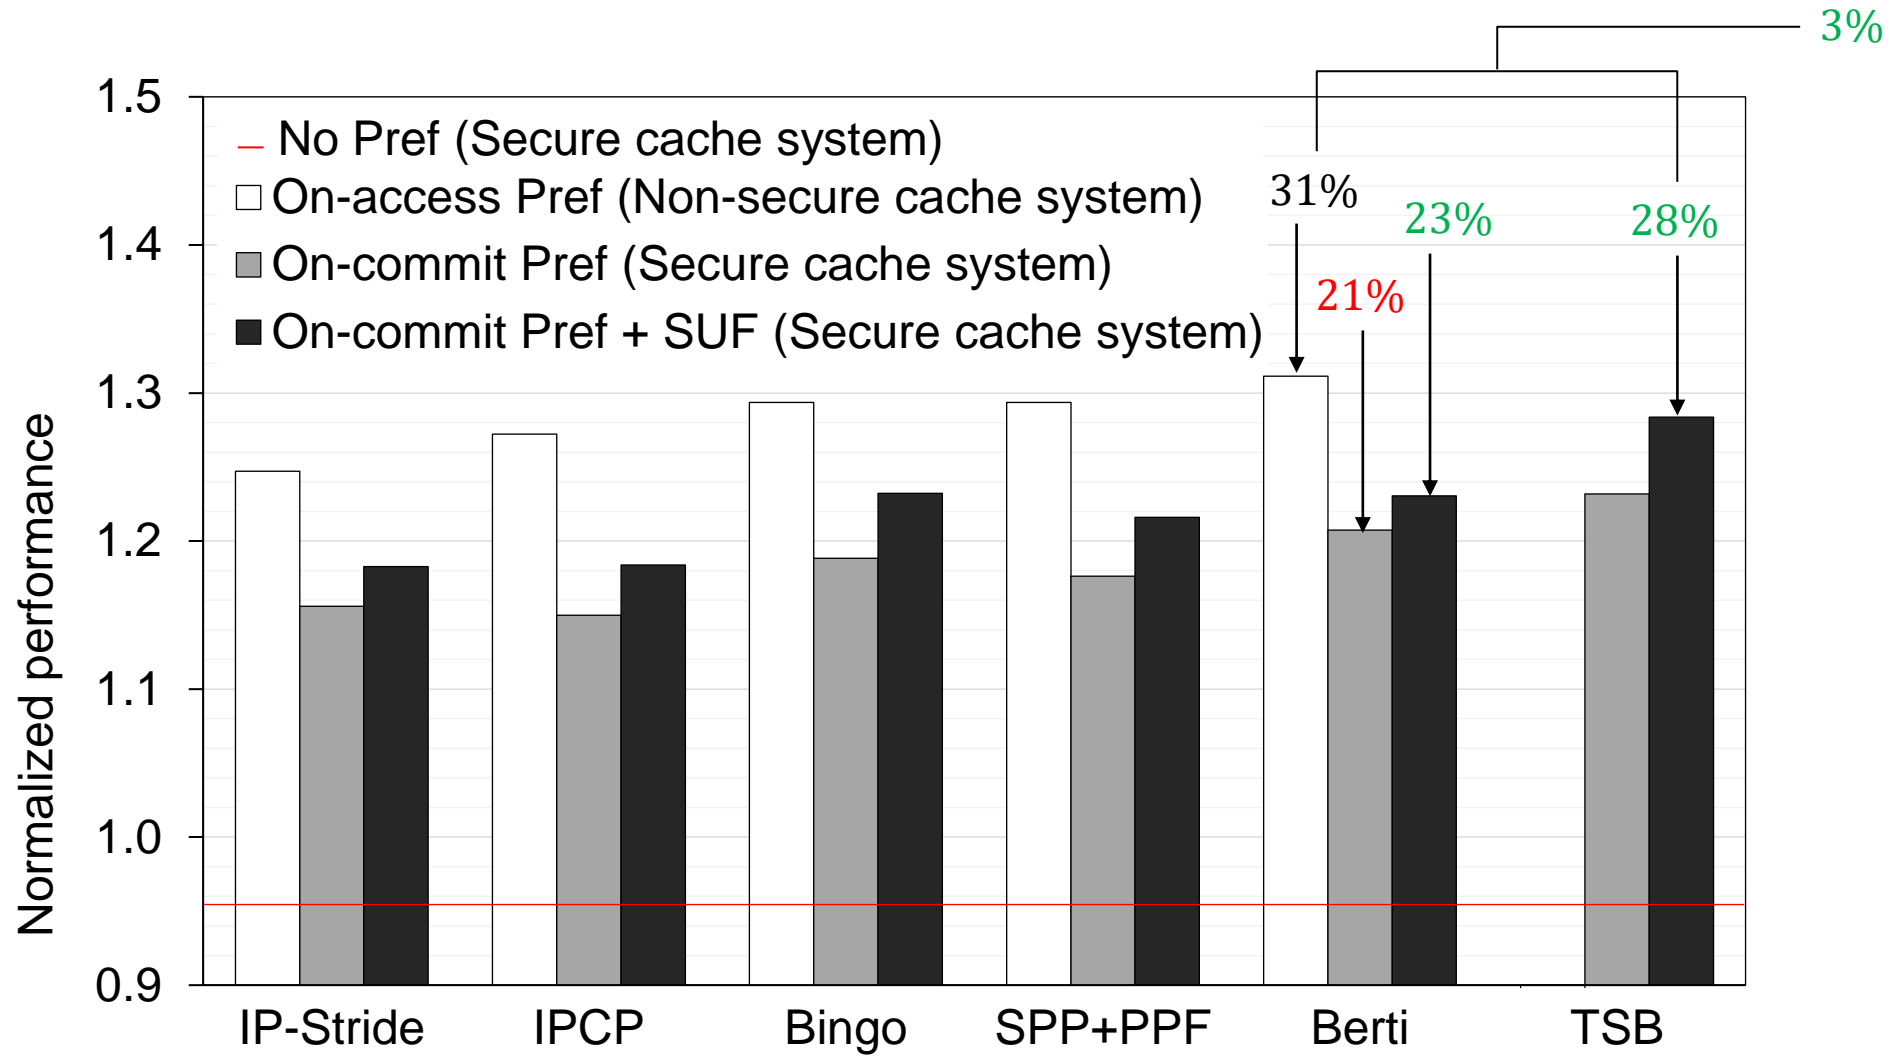
TSB: Hardware implementation



Evaluation setup


- Champsim(Intel Sunny Cove)
- Normalized to non-secure cache system without prefetching
- SPEC CPU 2017 & GAP benchmarks

Performance improvement




TSB improves performance by 6.3 % compared to On-commit Berti

Power reduction



SUF reduces dynamic energy consumption by 9.1%

Multi-core performance



TSB+SUF improves performance by 16.1% over baseline

Summary

- ▶ **Secure cache systems** hinders prefetcher performance
- ▶ Prefetchers are not built for **on-commit training/trigger**
- ▶ **SUF , TSB** improves prefetcher performance with security guarantees

Thanks and fingers crossed!

- ▶ Collaboration with **Prof. Alberto Ros and Agustin**, University of Murcia
- ▶ Work under submission to **ISCA 2024**, flagship forum in the field of computer architecture

Thank You

Background: Spectre^[S&P '19](variant 1)

```
if ( x < array1_size )  
    y = array2[ array1[ x ] * 512 ]
```



cached

uncached

11

E2

45

CD

- **x** , adversary controlled
- Adversary wants to read some **out of bounds** value of **array1**
- **array1_size** and **array2[]** not in cache
- Mistrains branch predictor

```
array2 [ 0 * 512 ]  
array2 [ 1 * 512 ]  
array2 [ 2 * 512 ]  
array2 [ 3 * 512 ]  
array2 [ 4 * 512 ]  
array2 [ 5 * 512 ]  
array2 [ 6 * 512 ]  
array2 [ 7 * 512 ]  
array2 [ 8 * 512 ]  
array2 [ 9 * 512 ]  
array2 [ 10 * 512 ]  
array2 [ 11 * 512 ]  
array2 [ 12 * 512 ]
```

Background: Spectre^[S&P '19] (variant 1)

```
if ( x < array1_size )  
    y = array2[ array1[ x ] * 512 ]
```



cached

uncached

Secret

11

E2

45

CD

- ▶ Attacker calls victim with out-of-bounds x
 - ▶ Wrongly predicts true and starts Speculative exec
 - ▶ Reads addr (array1 base + x) with out-of-bounds x
 - ▶ Read returns Secret byte = **11**
 - ▶ Brings array2 [**11** * 512] into cache
 - ▶ Branch resolves: instructions squashed
- ▶ Attacker times reads from array2[i * 512]
 - ▶ Read for i = **11** is fast, revealing secret data

```
array2 [ 0 * 512 ]  
array2 [ 1 * 512 ]  
array2 [ 2 * 512 ]  
array2 [ 3 * 512 ]  
array2 [ 4 * 512 ]  
array2 [ 5 * 512 ]  
array2 [ 6 * 512 ]  
array2 [ 7 * 512 ]  
array2 [ 8 * 512 ]  
array2 [ 9 * 512 ]  
array2 [ 10 * 512 ]  
array2 [ 11 * 512 ]  
array2 [ 12 * 512 ]
```