CASPER

GhostMinion: A Strictness-Ordered Cache System for Spectre Mitigation[1]

Sumon Nath
sumon@cse.iitb.ac.in

[1] Ainsworth, Sam. "GhostMinion: A strictness-ordered cache system for Spectre mitigation." In *MICRO-54: 54th Annual IEEE/ACM International Symposium on Microarchitecture*, pp. 592-606. 2021.

Recap- Spectre

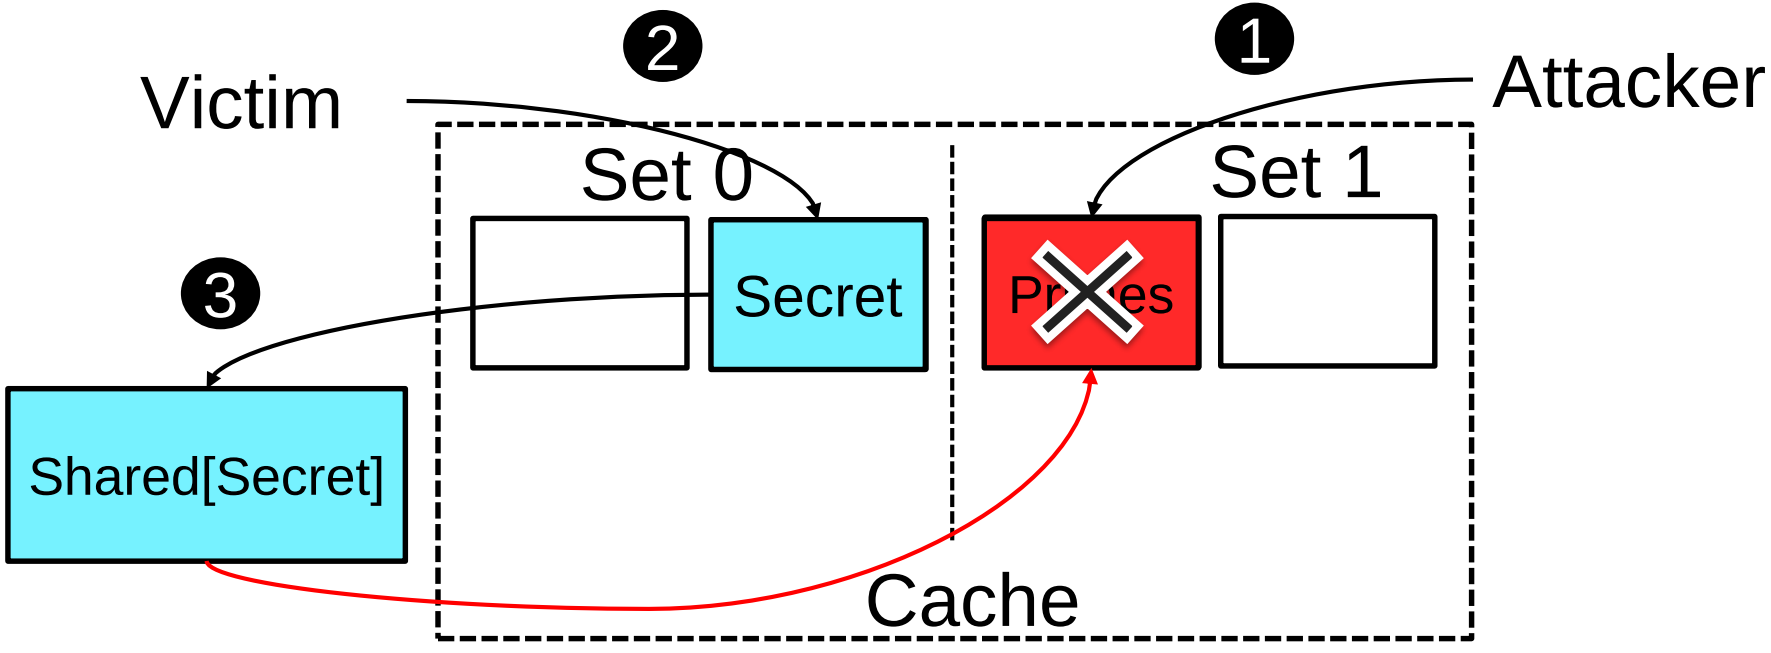

Problem: Access to speculative data across **domain boundaries**

Recap- MuonTrap

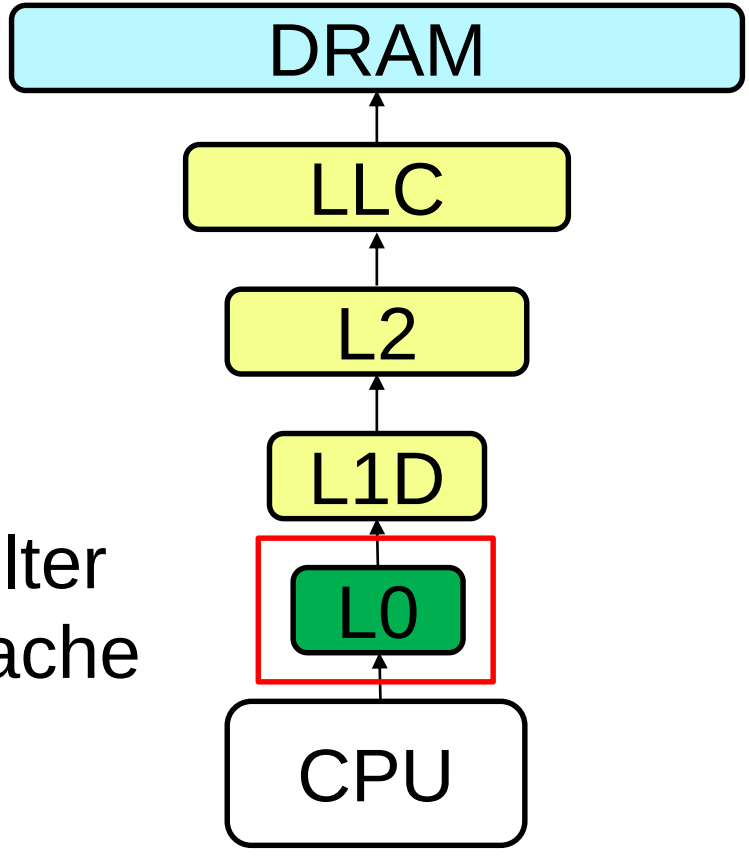

- Stores all speculative data
- Wiped on context switch
- Non-inclusive/non-exclusive
- Instruction cache, TLBs

Does this work always?

Attack scenario

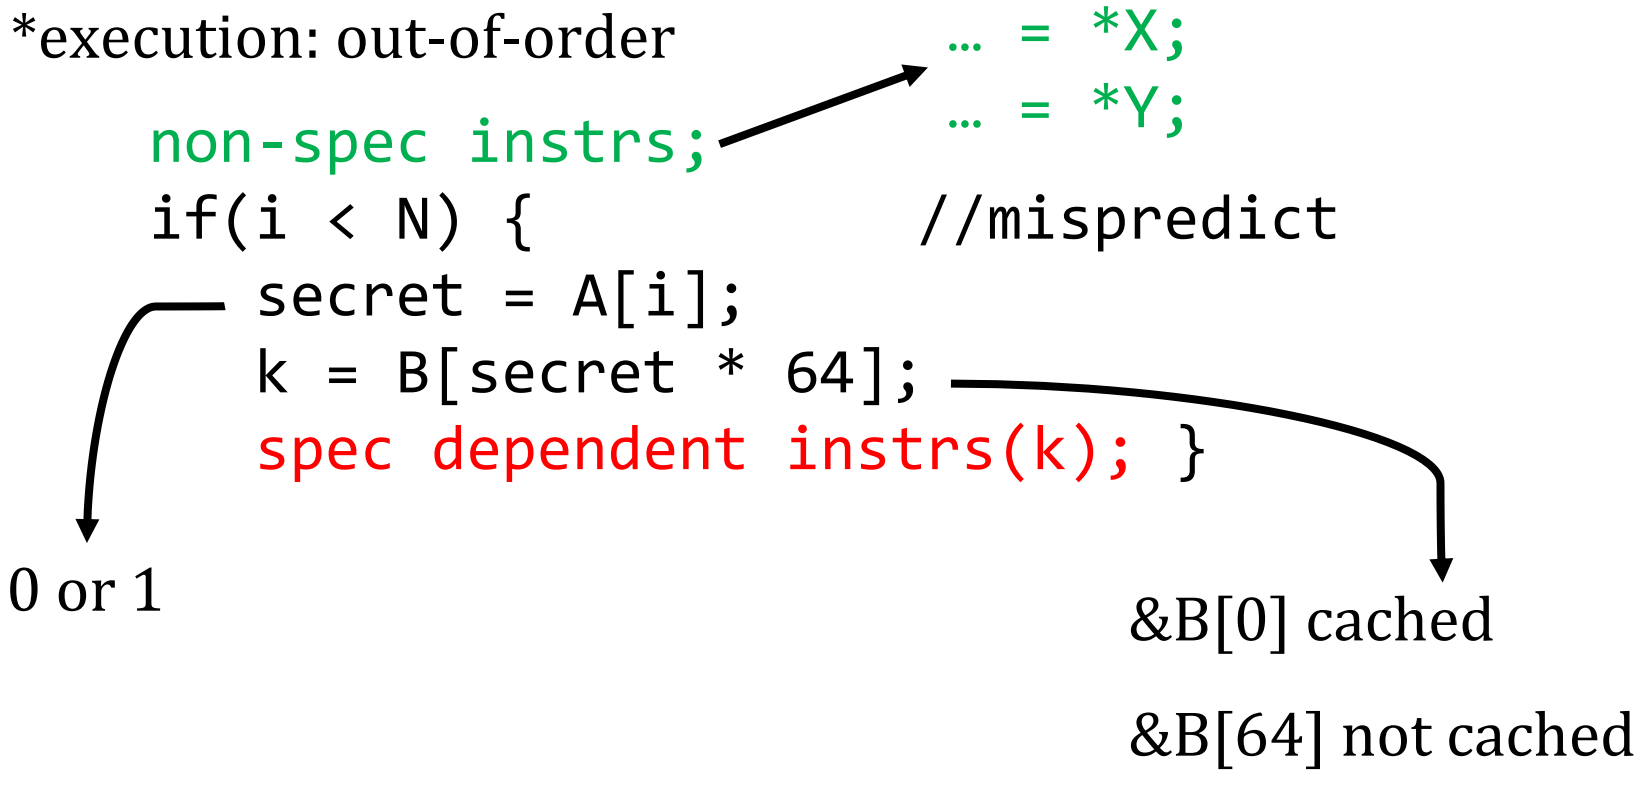

Execution unit contention

```
non-spec instrs;
if(i < N) {                    //mispredict
    secret = A[i];
    k = B[secret * 64];
spec dependent instrs(k); }
```

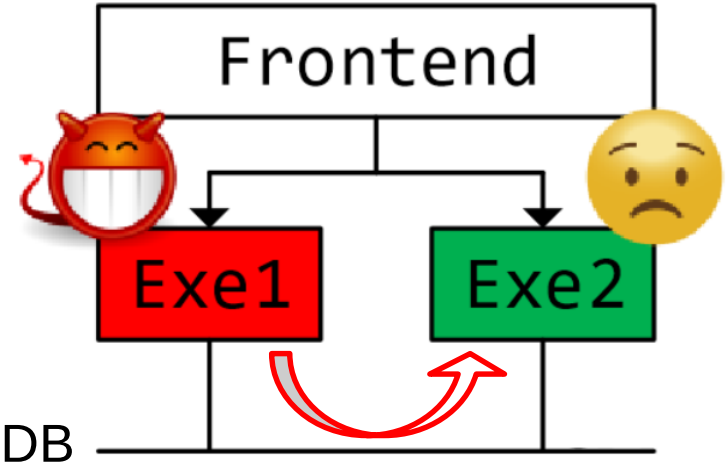

Attack scenario

```
non-spec instrs;
if(i < N) {                    //mispredict
    secret = A[i];
    k = B[secret * 64];
spec dependent instrs(k); }
```

Speculative Interference[2]

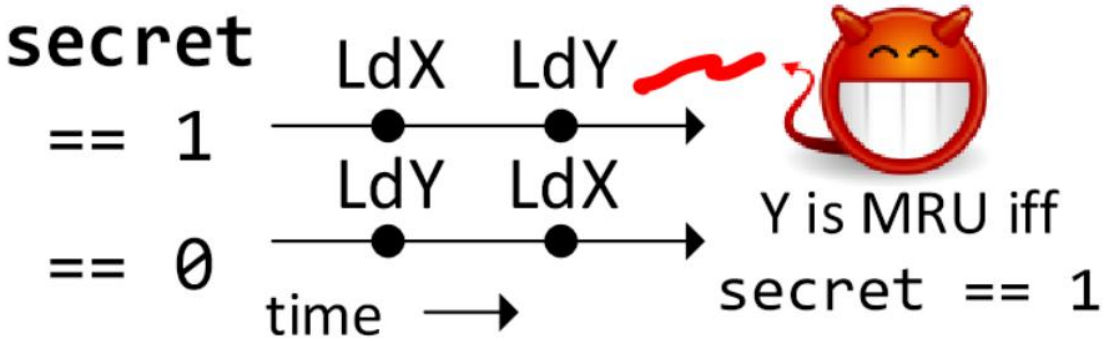

[1] Behnia, Mohammad, Prateek Sahu, Riccardo Paccagnella, Jiyong Yu, Zirui Neil Zhao, Xiang Zou, Thomas Unterluggauer et al. "Speculative interference attacks: Breaking invisible speculation schemes." In *Proceedings of the 26th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, pp. 1046-1060. 2021.

Does MuonTrap work?

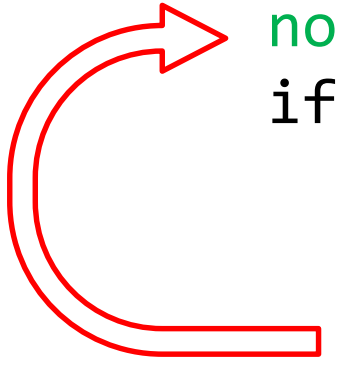

```
non-spec instrs;
if(i < N) {                    //mispredict
    secret = A[i];
    k = B[secret * 64];
    spec dependent instrs(k); }
```

Does not take care of backwards-in-time channels

No!

Till now…

- Spectre
- MuonTrap
- Speculative Interference
- Why MuonTrap does not work?
- Next: GhostMinion

What can be done?

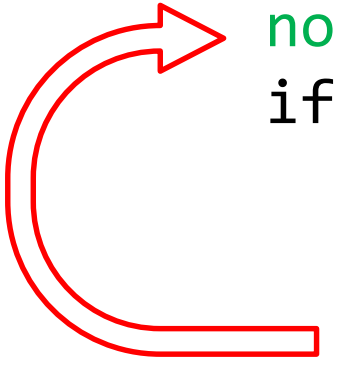

```
non-spec instrs;
if(i < N) {                    //mispredict
    secret = A[i];
    k = B[secret * 64];
    spec dependent instrs(k); }
```

We need to restrict backward-in-time channels

How? 🤔

Strictness ordering

x can impact timing of y, *iff*

$$\mathbf{commit}(y) \not\rightarrow \mathbf{commit}(x)$$

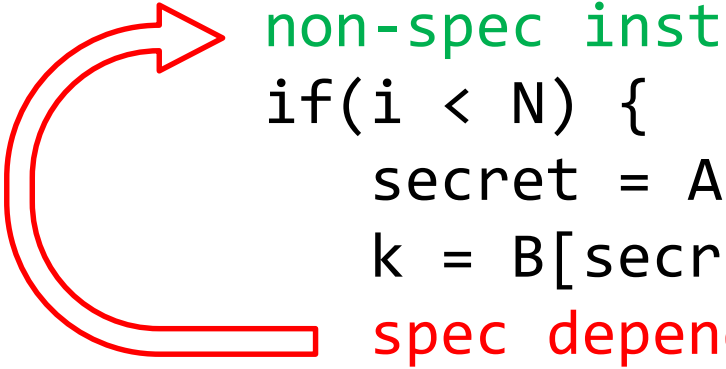

```
non-spec instrs;
if(i < N) {
    secret = A[i];
    k = B[secret * 64];
    spec dependent instrs(k); }
```

y ***commits***

x ***doesn't commit***

Temporal ordering

*strictness ordering hard to implement

$$x \textbf{ can impact timing of } y, \textit{iff}$$
$$\mathbf{commit}(x) \vee \mathbf{seq}(x, y)$$

Does it ensure strictness ordering?

yes! $\mathbf{seq}(x, y) \rightarrow \mathbf{commit}(y) \rightarrow \mathbf{commit}(x)$

Implementation

GhostMinion - cache system

Techniques:

- TimeGuarding
- Free-slotting
- LeapFrogging

*Applied to other microarchitecture structures

GhostMinion

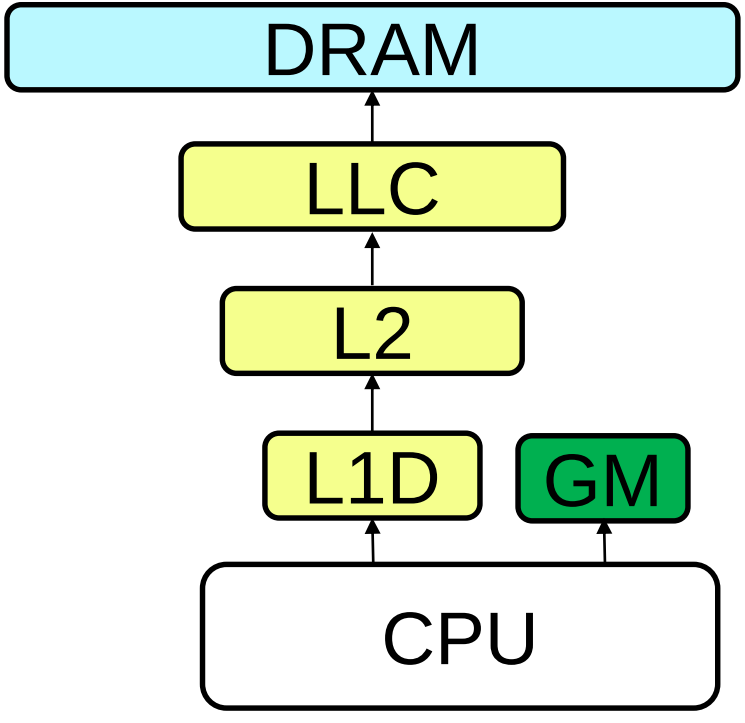

- Same as MuonTrap
- Accessed in parallel to L1D

TimeGuarding

Ensures temporal ordering

Read

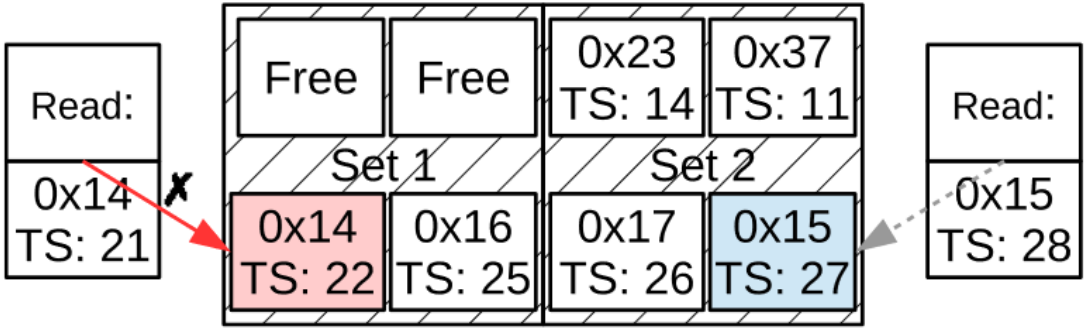

Fill

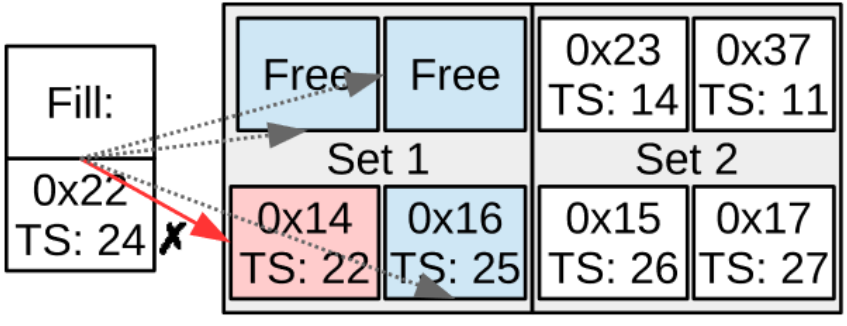

Free-slotting

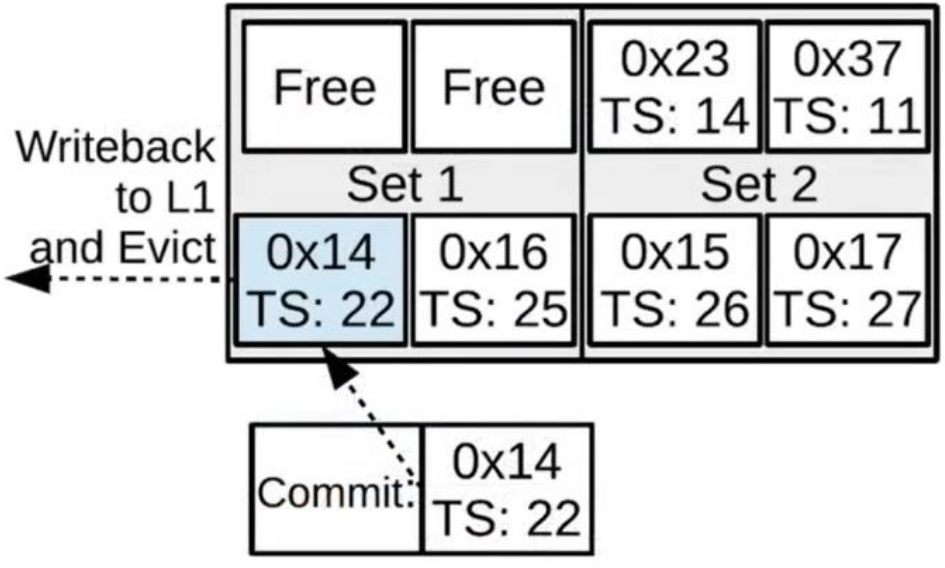

Avoids resource starvation

LeapFrogging

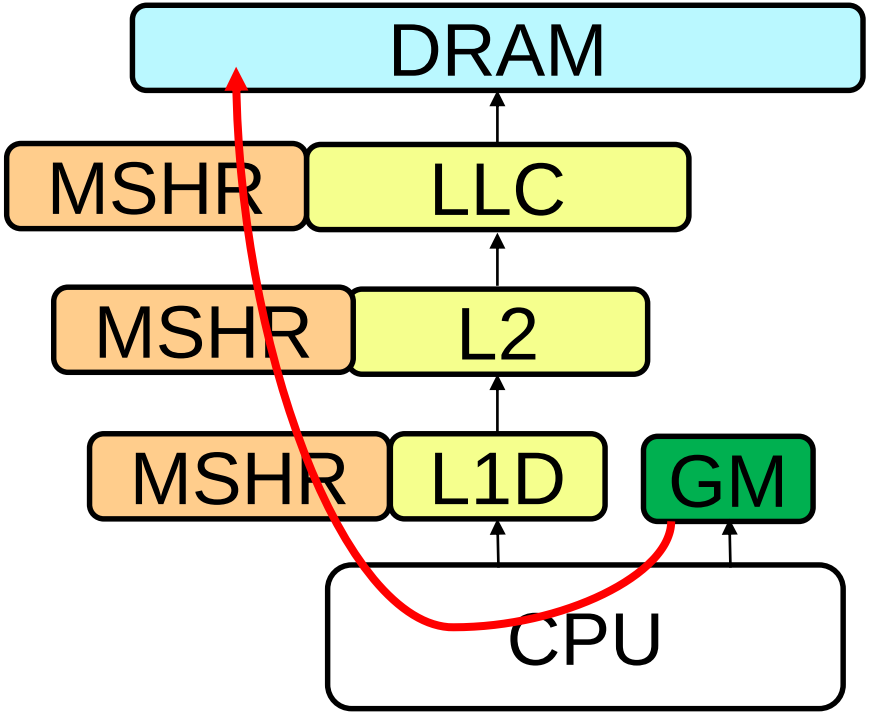

- MSHR:
  stores miss status
  enables parallel misses
- Non-inclusive/
  non-exclusive cache

Attack scenario

```
1 A = ...     // takes Z cycles
2 y = load(A)   // Interference Target
3 if (i < N): // mispred. taken (miss on N)
4     secret = load(&TargetArray[i]) // access
5     // Interference Gadget
6      x0    = load(&S[secret * 64 * 0])
7      x1    = load(&S[secret * 64 * 1])
8       ...
9      x_{M-1} = load(&S[secret * 64 * (M-1)])
```

Attack scenario

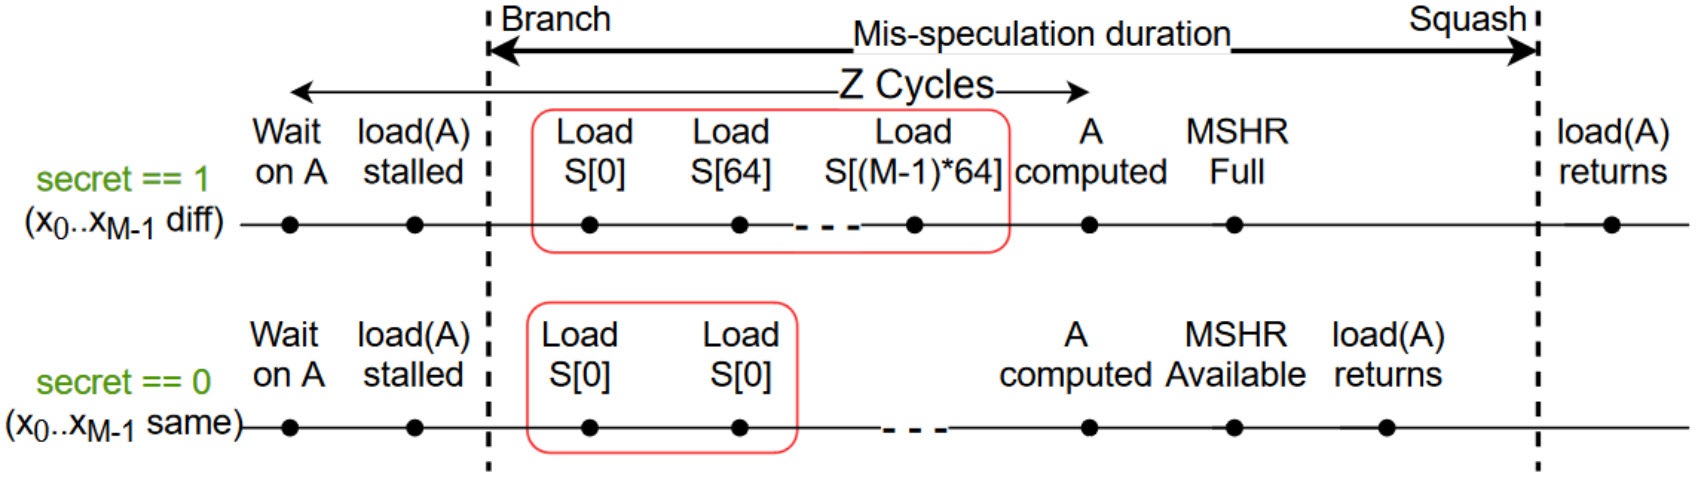

violates strictness ordering

LeapFrogging

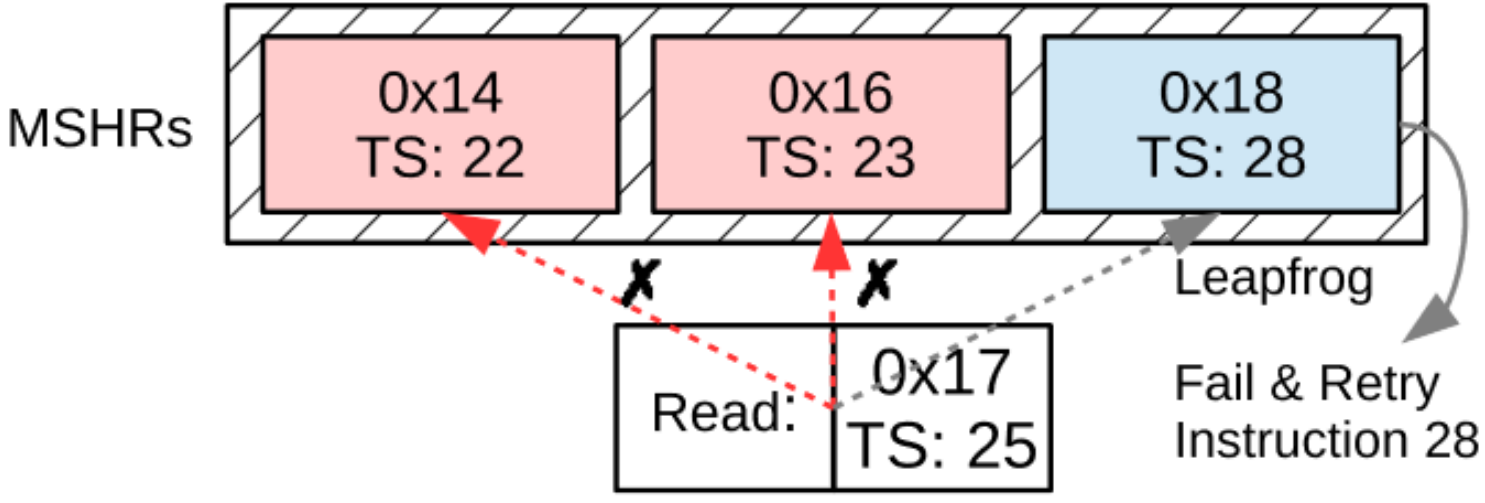

Older instruction can kick out MSHR entry of newer instruction

Performance evaluation: SPEC 2006

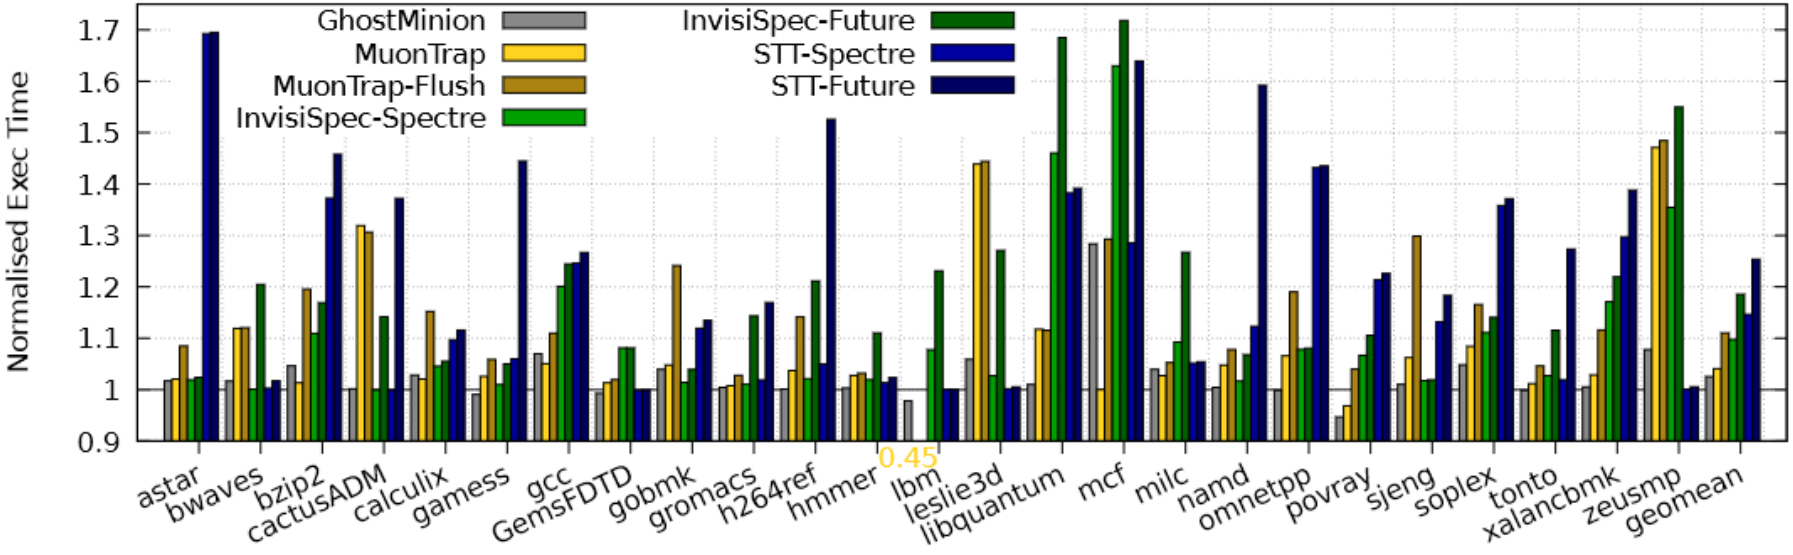

- GhostMinion 2.5% slowdown compared to 4% slowdown for MuonTrap

Conclusion

- MuonTrap was the first to solve Spectre with comparatively low performance overhead.
- GhostMinion proposes a precise framework to avoid Spectre and its different variants.

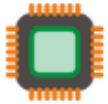

CASPER

Thank You!

Speculative Interference


```
1  z = ...      // takes Z cycles
2  A = f(z)     // takes F cycles
3  y = load(A)
4  B = g(z)     // takes G > F cycles
5  v = load(B)
6  if (i < N): // mispredict taken (miss on N)
7      secret = load(&TargetArray[i])
8      // Interference Gadget
9      x = load(&S[secret * 64]) // secret=1->hit, secret=0->miss
10     f'(x)
```

contention on a non-pipelined EU. Instruction sequences `f` and `f'` use the same non- pipelined EU. Instruction sequence g uses a different EU